

Gaussian Process Regression

A practical overview

Tzu-Heng Lin, 2014011054, W42

Department of Electronic Engineering, Tsinghua University, Beijing, China

lzhbrian@gmail.com

ABSTRACT

¹Gaussian Process Regression(GPR) is a powerful, nonparametric tool developed based on the Bayesian theory and the Statistical learning theory. Choosing the right *Mean Functions*, *Kernel Functions* as well as the *Likelihood Functions* and the *Inference Methods* have been critical to the performance of the model. However, these works are often hard and require much expertise & experience.

In this paper, we first give an introduction on the overall process of GPR. Subsequently, we give a precise explanation on some of the recent works which emphasize on the choice of the *Kernel Function*. In addition, we implement sufficient number of experiments to systematically analyze the performance of GPR with different *Likelihood Functions* and the *Inference Methods*. Our experiments are conducted on two interesting datasets. The best MSE we get from two experiments are 0.2112 & 0.0262.

We seek to provide an comprehensive practical overview in the field of Gaussian Process Regression.

1. INTRODUCTION

Machine learning has been a heated research topic these days. With this amazing tool, we are now capable of predicting the price of the stock price based on history, doing the classifying by just inputing the pixels of images.

Supervised learning is one of the most important sections for machine learning. And Regression is probably the core of supervised learning. By firstly inputing in the computer some of the training data, the machine would be able to learn the characteristics of the dataset and make predictions.

Gaussian Process Regression(GPR) is a supervised learning regression method which are getting increasingly welcome in both the research field and the industry. In a GPR, we take advantage of the flexibility and simplicity of a Gaussian Process and implement it into a regression problem.

There are still countless unsolved problems in the field of GPR. In this paper, we would comprehensively introduce the concept of GPR, including most of the notable works. Specifically, we focus on some marvellous kernel choosing works. We also implemented several experiments to quantitatively analyze the performance of different *Mean Functions*, *Likelihood Functions* and the *Inference Methods*.

¹Tzu-Heng Lin is currently an undergraduate student in the Department of Electronic Engineering, Tsinghua University. His research interests include Big Data Mining, Machine Learning, etc. For more information about me, please see [my personal website](#). Please feel free to contact me at any time via [email](#)

tions, *Likelihood Functions* and the *Inference Methods* We seek to provide a practical overview in the field of GPR.

The structure of this paper is as follows: In section 2, we provide a precise overview of a Gaussian Process Regression. In section 3, we briefly listed some of the notable progress on GPR. Section 4 describes some great recent work on the methods for auto-construction of the kernels, as well as an interesting additive kernel which can have a better performance when solving some specific cases. In section 5, we conduct two experiments and use them to compare the performance of different methods and algorithms. Conclusions are drawn in section 6.

2. GAUSSIAN PROCESS REGRESSION

2.1 Regression

Regression Regression is probably one of the most fundamental problems in a wide range of fields including *Statistics*, *Signal Processing* and *Machine Learning*, etc. A regression problem is usually formulated as follows: Given a training set $D = \{(\mathbf{x}_i, y_i) | i = 1, 2, \dots, n\}$, we assume that \mathbf{x}_i, y_i have the following relationship:

$$y_i = f(\mathbf{x}_i) + e \quad (1)$$

where e is the error noise. By finding such $f(\cdot)$, we can predict what a corresponding y^* is in some test case \mathbf{x}^* . Note that \mathbf{x} can either be a vector or a scalar.

Generalized Linear Model A widely used regression model is called Generalized Linear Model(GLM)[10], in which a regression function can be expressed as a linear combination:

$$f(\mathbf{x}) = \sum_{i=1}^M w_i \phi_i(\mathbf{x}) \quad (2)$$

where $\phi_i(x)$ is called the basis function. In a regular GLM analysis, we have to firstly determine what our basis functions we are going to use, and subsequently can we use the training dataset to derive the parameters in the basis functions and the coefficients in the regression function.

Mean Square Error We use a measurement called the Mean Square Error(MSE) to evaluate the performance of the regression function. It is defined as follows:

$$MSE = \frac{1}{m} \sum_{i=1}^m (f(\mathbf{x}^*) - y_i^*)^2 \quad (3)$$

where $f(x)$ represents the regression function. A smaller MSE represents a better regression function on a test set.

2.2 Gaussian Process Regression

Gaussian Process A Gaussian Process(GP) is any distribution over functions such that any finite set of function values $f(x_1), f(x_2), \dots, f(x_N)$ have a joint Gaussian distribution[15]. It can usually be represented as

$$N\{E[f(x)], Cov[f(x), f(x')]\} \quad (4)$$

where $E[f(x)]$ refers to its *Mean Function*, and $Cov[f(x), f(x')]$ refers to its *Covariance Function*.

Gaussian Process Regression Gauss Process Regression(GPR)[15] is a popular regression method these years. The key of this method is to model the regression function $\{f(\mathbf{x})|\mathbf{x} \in S\}$ as a GP

$$f(x) \leftarrow N\{m(\mathbf{x}), K(\mathbf{x}, \mathbf{x}')\} \quad (5)$$

where $m(\mathbf{x})$ is the *Mean function* and $K(\mathbf{x}, \mathbf{x}')$ is the *Kernel Function*.

In a GPR, we don't have to derive the exact form of the regression function, we just need to determine the form of the above two functions. As introduced in GP, the *Kernel Function* $K(\mathbf{x}, \mathbf{x}')$ is actually the covariance between $f(\mathbf{x})$ and $f(\mathbf{x}')$, and if it is a zero-mean GP, the covariance turns into correlation. So a *Kernel Function* represents the relationship between $f(\mathbf{x})$ and $f(\mathbf{x}')$.

By calculating the posterior probability of the desired $f(\mathbf{x}^*)$, i.e. $p(f(\mathbf{x}^*)|\mathbf{x}^*, D)$, we can derive the mean value along with the standard deviation of this estimation. On the other hand, we must noted that calculating the posterior probability will become an intractable work when we have a high-dimensional dataset. It is a need that we introduce some inference method to estimate this work.

To summarize, choosing a suitable *Mean Functions*, *Kernel Functions* as well as the *Likelihood Functions* and the *Inference Methods* is the key of a GPR model. Rasmussen's *Gaussian Processes for Machine Learning*[15] has implemented some marvellous Matlab/Octave code of GPR, it is available on his website² known as **GPML**.

3. RELATED WORK

Gaussian Process Regression Rasmussen[15] and Williams[18] have a wonderful fundamental work on the introduction of Gaussian Process Regression.

Kernel Choice Duvenaud et.al introduce a GP model of functions which are additive[4]. Along with Lloyd et.al, they introduce a way to automatic construct some desired kernels to better fit the dataset[3, 2]. They also figure out an auto way to produce the Natural-Language description of a data structure[8]. In this paper, we will give a precise introduction on the above work.

Inference Methods Some popular works on inference methods include MCMC[6], Expectation Propagation[11], Variational Bayes[14, 13], Leave-One-Out[5], and Laplace Approximation[17], etc. Due to the limited time and space, we won't go very deep into these inference methods.

4. KERNEL CHOICE

Choosing an appropriate *Kernel Function* in a regression problem has always been a nightmare for researchers and analysts. In this section, from the practical meaning of a *Kernel Function*, we introduce an automatic way to construct a desired composite form of *Kernel*. We also introduce a special *Additive Kernel* used specifically for an Additive Gaussian Process, which is going to be a useful method in our experiment (Section 5). This section is based on the work of Duvenaud and Lloyd et. al [3, 8, 2, 4].

4.1 Kernels express structures

Base Kernels As we all know, different kernels can represent different kinds of data. In Figure 1, we show 4 different kinds of base kernels, each represents a specific kind of data characteristics (Table 1). Each function of the kernel is shown in Equation 6

$$\begin{cases} k_{LIN}(x, x') = \sigma_b^2 + \sigma_v^2(x - l)(x' - l) \\ k_{SE}(x, x') = \sigma^2 \exp(-\frac{(x - x')^2}{2l^2}) \\ k_{PER}(x, x') = \sigma^2 \exp(-\frac{2\sin^2(\pi(x - x')/p)}{l^2}) \\ k_{RQ}(x, x') = \sigma^2(1 + \frac{(x - x')^2}{2\alpha l^2})^{-\alpha} \end{cases} \quad (6)$$




Figure 1: Base Kernels

Kernel	Data Structures
Linear(LIN)	linear functions
Squared exponential(SE)	local variation
Periodic(PER)	repeating structure
Rational quadratic(RQ)	multi-scale variation

Table 1: Different kinds of kernels and its represented data structures.

Compositional Kernels When the data structure we are dealing with is not contained in any of the above base kernel, we have to made one to fit the targeted data characteristics. A possible and probable approach of making a

²Available at <http://www.gaussianprocess.org/gpml/code/>

customized kernel is to do some addition and multiplication to the base kernels.

$$k_{k_a+k_b}(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}, \mathbf{x}') + k_b(\mathbf{x}, \mathbf{x}') \quad (7)$$

$$k_{k_a \times k_b}(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}, \mathbf{x}') \times k_b(\mathbf{x}, \mathbf{x}') \quad (8)$$

By addition and multiplication, we are able to construct some very complicated data structure. Below, we show some examples of structures by compositional kernels (Figure 2) and the data structures they are capable of representing (Table 2).




Figure 2: Some examples of the compositional kernels

Kernel	Data Structures
LIN+ PER	periodic plus trend
SE + PER	periodic plus noise
SE x PER	locally periodic
LIN x LIN	quadratic functions

Table 2: Some examples of the compositional kernels

4.2 Automatic Construction

As we can see in the above analysis, choice of the form of a *Kernel Function* can be critical in a GPR. However, this process was used to be a privileged work for experts since it requires a whole lot of experience and expertise. In [3, 8, 2], Duvenaud and Lloyd et.al describe a tractable method to let a computer be capable of doing this work. This procedure is summarized as follows:

Kernel Family We consider the *base kernels* as shown in Figure 1, any algebraic expression using these kernels as a combination of + and \times could be our target. And any of our target with the concatenation of the parameters forms a kernel family.

Scoring a Kernel Family We must find a way to evaluate a kernel family. At here, we use the marginal likelihood[1] as our criterion, and use the Bayesian Information Criterion(BIC)[16] to approximate the integration over kernel parameters.

Search over structures Last, by using a greedy search: At each stage, expand the current kernel by all possible *op-*

Algorithm 1 Automatic kernel construction algorithm

Require: Required search depth N
Initialize kernel S
for $t = 1 \rightarrow N$ **do**
 $CurScore \leftarrow 0$
 for Possible operators S' **do**
 if $Score(S) > CurScore$ **then**
 $S \leftarrow S'$
 end if
 end for
end for
return kernel S

erators and choose the highest scoring one. The number of working stage is defined by user. Possible operators are listed as follows:

1. Any expression S can be replaced by $S + B$
2. Any expression S can be replaced by $S \times B$
3. Any base kernel B can be replaced by B'

where B and B' represent any base kernel family. Note that, in the work [8], Lloyd et. al also take account of the changepoint operator CP , the changewindow operator CW , and some empirical operators. They also add some base kernels to improve the algorithm. Due to limited time and space, we only introduce the most important part here.

The implemented Matlab and Python code by Duvenaud and Lloyd et.al can be found on github³.

4.3 Additive Gaussian Processes

Additive Gaussian Process In our daily life, many natural functions result from some unknown interactions of some base kernels. For example, the price of a house may be jointly dependent on some features such as the size, the location and the materials, etc. In this method, we describe a way that can automatically discover the hidden interactions between the dimensions.

Additive Kernels We now give the precise definition of additive kernels. As shown in [4], we define the 1st, 2nd, nth order additive kernel as:

$$k_{add_1}(\mathbf{x}, \mathbf{x}') = \sigma_1^2 \sum_{i=1}^D k_i(x_i, x'_i) \quad (9)$$

$$k_{add_2}(\mathbf{x}, \mathbf{x}') = \sigma_2^2 \sum_{i=1}^D \sum_{j=i+1}^D k_i(x_i, x'_i) k_j(x_j, x'_j) \quad (10)$$

$$k_{add_n}(\mathbf{x}, \mathbf{x}') = \sigma_n^2 \sum_{1 \leq i_1 < i_2 < \dots < i_n \leq D} \left(\prod_{d=1}^n k_{i_d}(x_{i_d}, x'_{i_d}) \right) \quad (11)$$

$$(12)$$

where $k_i(x_i, x'_i)$ is the *base kernel* we assigned at first for each dimension $i \in \{1, 2, \dots, D\}$. A full additive kernel is a sum of all orders' additive kernels, as in equation 13:

$$K_{add_{full}}(\mathbf{x}, \mathbf{x}') = \sum_{n=1}^D k_{add_n}(\mathbf{x}, \mathbf{x}') \quad (13)$$

³Available at github.com/jamesrobertlloyd/gpss-research and github.com/jamesrobertlloyd/gp-structure-search.

Practice In the real practice, the only uncertain thing should be decided by us is the *base kernels*. By using this *Additive Kernel* method, we just need to input a sum of $1 : R$ th order additive kernel(A full additive kernel) and it will automatically optimize the contribution of each order.

Special Case We also note that if all of the contribution of the additive GP comes from the 1st order, it turns into a Generalized Additive Models(GAM)[7], and if all of the contribution comes from the D th order, it turns into an SE.

5. EXPERIMENTS

5.1 First Experiment

Dataset Our first experiment is conducted on a one-dimensional data as shown in Figure 3. In this experiment, we have a 543-point training set ranging from $x = -32$ to $x = 14$, and we want to predict some 67-point test set ranging from $x = 14$ to $x = 21$.




Figure 3: A one-dimensional dataset

Auto Kernel Construction To better do this work, we use the Auto Construction Kernel method described in Section 4.2. By doing a lengthy greedy kernel construction procedure, we derive a compositional kernel:

$$LIN \times SE + SE \times (PER + RQ) \quad (14)$$




Figure 4: Prediction of the 1st dataset

Result And by using the *Linear*, Mean Function *Gaussian* Likelihood function and the *Exact* posterior probability. (Note that we have also optimized the hyperparameters in all experiments.) We predict the results as in Figure 4. We

could achieve an MSE of **0.2112** using this method. Given that the output data ranges from -60 to 30 , we believe that we have a satisfying result.

5.2 Second Experiment

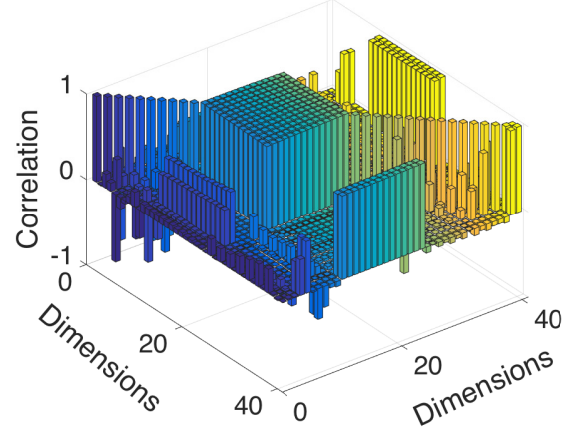


Figure 5: Ailerons dataset, Correlation between each dimension of \mathbf{x} . We can see that in the dimension set $\{11:24, 39,40\}$, variables are highly correlated with each other.

Dataset The second dataset called *Ailerons*⁴ is from the *Experiments of Rui Camacho*(rcamacho@garfield.fe.up.pt), it addresses a control problem of F16 aircraft. The goal is to predict the control action on the ailerons of the aircraft. This dataset contains 40 attributes with a 10,000-point training set and a 3,750-point test set.

Preprocessing Given a such high-dimensional dataset, we decide to first analyze the relationship between each dimension. As shown in Figure 5, we can see that actually 16 dimensions of the 40-dimensional \mathbf{x} are highly correlated to each other (correlation > 0.96). So reducing the dimension of this dataset is a probable way to reduce the complexity of the algorithm. Thus, we do this procedure by ignoring the $\{12:24, 39,40\}$ th dimensions, which are all highly correlated with the 11th dimension. And we consequently derive a 25-dimensional set with the $\{1:11, 25:38\}$ th dimensions.

As follows, we are going to use two respective methods, the first one is through the auto construction of kernels, the second one is through the characteristics of an additive kernel as described in Section 4.3.

5.2.1 Auto Kernel Construction

By conducting a greedy kernel search as in the previous experiment, we get a desired kernel function: *SE*. We hereby use this kernel and test its performance on different *Likelihood Functions* and *Inference Methods*. We use two different *Likelihood Functions*: Gaussian and Laplace, and three *Inference Methods*: Laplace, Variational Bayesian(VB) and Leave-One-Out(L00). The results are shown in Table 3

Note that due to the large dataset and limited computing facilities, we only randomly sampled 1,000 points as the

⁴Available at www.dcc.fc.up.pt/~ltorgo/Regression/aileron

training set, and optimize the hyperparameters within 500 function evaluations. We set the *Mean Function* as a Constant. We also use the *Automatic Relevance Determination(ARD)*[9, 12], when doing the estimation.

MSE	Exact	Laplace	VB	L00
Gaussian	0.0263	0.0262	0.0262	0.0276
Laplace	N/A	0.9266	0.1680	0.0277
Time(s)	Exact	Laplace	VB	L00
Gaussian	106.28	148.56	457.26	170.48
Laplace	N/A	3.76	693.64	185.32

Table 3: MSE and Time(second) using different kinds of *Likelihood Functions* and *Inference Methods*, using kernel SE. (Different rows represents different *Likelihood Functions* and different columns represents different *Inference Methods*).

5.2.2 Additive Kernel

Dataset Structure Given that it is a high-dimensional dataset, we decide to try using the *Additive Kernel* method to discover the hidden relationship between the variables.

Additive Kernel Since we have a 25-dimension(reduced) dataset, we set the kernel functions as:

$$\begin{aligned}
K_{add} &= \sum_{r=1}^R k_{add_r}(\mathbf{x}, \mathbf{x}') \\
&= \sum_{r=1}^R \{ \sigma_r^2 \sum_{1 \leq i_1 < i_2 < \dots < i_r \leq D} (\prod_{d=1}^r k_{i_d}(x_{i_d}, x'_{i_d})) \}
\end{aligned} \tag{15}$$

where

$$\begin{aligned}
D &= 25 \\
R &= 3 \\
k_{i_d}(x_{i_d}, x'_{i_d}) &= k_{SE}(x_{i_d}, x'_{i_d}) \\
&= \exp\left(-\frac{(x_{i_d} - x'_{i_d})^2}{2l_{i_d}^2}\right)
\end{aligned} \tag{16}$$

Due to the complex kernel, the optimization of the hyperparameter becomes too lengthy, so we only use a sampled 200-point train set, and with a 100 time limited function evaluation. We also set $R = 3$ because a large R is really far from our Mac’s capability to compute. We use the same analysis method as in section 5.2.1, the result is shown in Table 4. Contribution of each order are shown in Table 5.

MSE	Exact	Laplace	VB	L00
Gaussian	0.0359	0.0357	0.0360	0.0356
Laplace	N/A	∞	0.0352	0.0362
Time(s)	Exact	Laplace	VB	L00
Gaussian	44.45	47.83	116.12	46.14
Laplace	N/A	6.35	483.57	46.31

Table 4: MSE and Time(second) using different kinds of *Likelihood Functions* and *Inference Methods*, using an additive kernel with SE.

No.	1	2	3
Contribution	8.16	17.32	74.52

Table 5: Contribution of each order additive kernel with the best performance(Likelihood:Laplace, Inference:Variational Bayes). (Normalized to sum to 100.)

5.3 Analysis

Auto Kernel Construction We only use a 1,000 point train set, but achieve a **0.0262 Mean Square Error**, which is very impressive. However, to our surprise, although using most inference methods won’t disturb the results, they often cost more time(ridiculus). We guess its because the characteristic of this dataset or maybe the Kernel we have chosen is too simple.

Additive Kernel We did not achieve a better result than using the previous method, but note that we only use a 200 point dataset here and limit to only 100 function evaluations when doing the hyperparameter optimization. We also only set R to 3, which is significantly lower than the suggested 25. We strongly believe that if we have a better equipment and have more time, we could have a better result. And note that with such little training data and time of optimization, we still can achieve an **MSE of 0.0352**.

6. CONCLUSION

In this paper, we discuss about the Gaussian Process Regression. We comprehensively introduce the concept of Gaussian Process Regression. Specifically, we precisely explain the recent works for kernel auto construction. We also conduct two experiments, in which we systematically compare the performance of two Likelihood Functions and three Inference Methods. **The best MSE we get from two experiments are 0.2112 & 0.0262.**

As future work, we would like to join more information of GPR to this paper such as the details of how a hyperparameter is derived and some deeper discussion about the Inference methods; and if could, propose some improvement to the algorithms available.

7. ACKNOWLEDGEMENT

I would like to thank Yuanxin Zhang, XueChao Wang, for the discussion with me on the algorithms. Without them, I wouldn’t have the possibility to accomplish this work in such a short time. This paper is a project of Stochastic Process Course in Tsinghua University, taught by Prof. Zhijian Ou.

Source Code and Dataset Source Code to perform all experiments, along with the dataset in this paper can be found in my github repository⁵. Due to my limited knowledge, there might be some mistakes and flaws in this paper as well as in the code, please do not hesitate to contact and correct me.

8. REFERENCES

- [1] A. Blumer, A. Ehrenfeucht, D. Haussler, and M. K. Warmuth. Occam’s razor. *Information processing letters*, 24(6):377–380, 1987.

⁵Available at <http://github.com/lzhhbrian/gpr>

- [2] D. Duvenaud. *Automatic model construction with Gaussian processes*. PhD thesis, University of Cambridge, 2014.
- [3] D. K. Duvenaud, J. R. Lloyd, R. B. Grosse, J. B. Tenenbaum, and Z. Ghahramani. Structure discovery in nonparametric regression through compositional kernel search. In *ICML (3)*, pages 1166–1174, 2013.
- [4] D. K. Duvenaud, H. Nickisch, and C. E. Rasmussen. Additive gaussian processes. In *Advances in neural information processing systems*, pages 226–234, 2011.
- [5] K. Fukunaga and D. M. Hummels. Leave-one-out procedures for nonparametric error estimates. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 11(4):421–423, 1989.
- [6] D. Gamerman. Sampling from the posterior distribution in generalized linear mixed models. *Statistics and Computing*, 7(1):57–68, 1997.
- [7] T. J. Hastie and R. J. Tibshirani. *Generalized additive models*, volume 43. CRC press, 1990.
- [8] J. R. Lloyd, D. Duvenaud, R. Grosse, J. B. Tenenbaum, and Z. Ghahramani. Automatic construction and natural-language description of nonparametric regression models. *arXiv preprint arXiv:1402.4304*, 2014.
- [9] D. J. MacKay. Bayesian methods for backpropagation networks. In *Models of neural networks III*, pages 211–254. Springer, 1996.
- [10] P. McCullagh. Generalized linear models. *European Journal of Operational Research*, 16(3):285–292, 1984.
- [11] T. P. Minka. Expectation propagation for approximate bayesian inference. In *Proceedings of the Seventeenth conference on Uncertainty in artificial intelligence*, pages 362–369. Morgan Kaufmann Publishers Inc., 2001.
- [12] R. M. Neal. *Bayesian learning for neural networks*, volume 118. Springer Science & Business Media, 2012.
- [13] H. Nickisch and M. W. Seeger. Convex variational bayesian inference for large scale generalized linear models. In *Proceedings of the 26th Annual International Conference on Machine Learning*, pages 761–768. ACM, 2009.
- [14] J. Palmer, D. Wipf, K. Kreutz-Delgado, and B. Rao. Variational em algorithms for non-gaussian latent variable models. *Advances in neural information processing systems*, 18:1059, 2006.
- [15] C. E. Rasmussen. Gaussian processes for machine learning. 2006.
- [16] G. Schwarz et al. Estimating the dimension of a model. *The annals of statistics*, 6(2):461–464, 1978.
- [17] L. Tierney and J. B. Kadane. Accurate approximations for posterior moments and marginal densities. *Journal of the american statistical association*, 81(393):82–86, 1986.
- [18] C. K. Williams. Prediction with gaussian processes: From linear regression to linear prediction and beyond. In *Learning in graphical models*, pages 599–621. Springer, 1998.